There is a pot of gold, you just don't see it: Applying Rainbow to partially observable environments

David Slayback
Northeastern University College of Computer and Information Science
440 Huntington Ave, No. 202
Boston, Massachusetts 02115

Abstract

Recent breakthroughs in deep reinforcement learning have led to AI solutions to many previously untenable games and environments. ViZDoom, an interface to the popular first person shooter (FPS) DOOM, is a particularly informative environment because of the partial observability of the state and potential complexity of the actions. This paper details an attempt to apply Rainbow, a reinforcement learning (RL) algorithm combining several improvements to Deep Q-Networks(DQN), to different scenarios within the ViZDoom environment. Preliminary results reveal the difficulty of the environment and offer insights into future directions for improving agent performance.

Introduction

In the past decade, with new access to data and computational power, artificial intelligence (AI) researchers have made enormous strides. Deep learning (DL), a type of machine learning once considered computationally intractable, has revolutionized our approach to many problems. Deep neural networks (DNNs) provide a generalizable alternative to previous domain-specific solutions.

Background

Reinforcement learning (RL) is a method by which some *agent* can learn to act in an unknown *environment* to maximize some goal expressed via *rewards*. RL is unique in that the agent does not need to be told *how* to act; rather, it learns through trial-and-error by encountering rewards as it acts and attempting to choose actions which maximize future rewards.

RL Terminology

The interaction of the agent and the environment is formalized as a Markov Decision Process (MDP). An MDP is defined by a tuple $\langle S, A, T, R, \gamma \rangle$ where:

- S is the set of possible states of the environment
- A is the set of possible actions of the agent
- T is the transition function $T(s, a, s')$ giving the probability of reaching each state s' from state s when taking action t
- R is the reward function $R(s, a)$ giving the reward of taking action a in state s
- γ is a discount factor applied to future rewards

At each timestep $t \in \tau$ (where τ is the end time), the agent takes an action a_t from start state s_t. The state changes to s_{t+1}, and the agent receives an immediate reward r_{t+1}. The agent chooses it's action based on a *policy* $\pi(a_t|s_t)$ which gives the probability of taking action a_t in state s_t. The objective of reinforcement learning is to achieve an optimal policy π^* which maximizes the reward the agent accumulates in the environment (i.e., the *return*, $G_t = \sum_{n=0}^{\infty} \gamma_t^k R_{t+k+1}$).

In value-based reinforcement learning, the agent learns π by determining the *Q-value* of its actions, where $Q^\pi(s, a)$ is the expected return of taking action a in state s if the agent chooses future actions based on policy π. It then derives a policy that takes the actions with the highest value, updating values as it encounters rewards. In *TD-learning*, the q-value is updated after each action, minimizing the difference between the *TD-target* (i.e., the currently estimated value of the state) and the combination of immediate reward and next state value seen (known as *TD-error*).

Large state spaces

One problem with this formalism is computational. In simple environments with few states and actions, an agent can store and calculate each $Q(s, a)$. In more complex environments, however, if the agent is even able to store all possible $Q(s, a)$, the number of actions it must take to accurately calculate them increases exponentially.

This issue can be solved by replacing the tables with DNNs, allowing the agent to learn reduced approximations of policies and values. Specifically, in the original DQN architecture* used for Atari games, the agent has a network that computes the Q-value for any input combination of pixels from the game frames. This network is parameterized by weights θ ($\widehat{\theta}$ being a periodic copy of the weights), such that the loss can be defined as:

$$(R_{t+1} + \gamma_{t+1} max_{a'} q_{\widehat{\theta}}(s_{t+1}, a') - q_\theta(s_t, a_t) \qquad (1)$$

Training an accurate DQN is data-intensive, but can leverage a *replay buffer* to reuse the agent's experiences and learn *off-policy* (i.e., without taking actions in the environment). This buffer of transitions $(s_t, a_t, r_{t+1}, \gamma_{t+1}, s_{t+1})$ can be sampled in random batches to provide additional training at only the cost of computational time.

Partial Observability

Another problem with the original formalism is sensory. In most environments, it's unrealistic for the agent to have access to the full state of the environment. Rather, the agent receives *observations* from its sensors in the form of camera images, communications, or other modalities. An extended formalism to capture this is known as a Partially Observable MDP (POMDP). A POMDP is an extended MDP defined by a tuple $\langle S, A, T, R, \Omega, O, \gamma \rangle$, where:

- Ω is the set of possible observations o
- O is the conditional probability function $O(o|s', a)$ specifying the probability of the agent observing o after taking action a and reaching environment state s'

At each timestep $t \in \tau$, the environment is in state s_t, the agent takes action a_t and receives observation $o_t + 1$ and reward $o_t + 1$. The environmental state may change, but the agent doesn't know. Rather, the agent determines a policy either directly from observation $\pi(a_t|o_t)$ or by developing a belief about the state s from its observations ($\pi(a_t|b_t)$ where b_t is a *belief* about the current state).

However, a single observation may not be sufficient to determine an optimal action or belief, so these are often conditioned on some history of actions, observations, and rewards h. Maintaining such a history is itself a problem, since it is difficult to know which experiences are relevant and impossible to store all of them.

Deep Recurrent Q-Networks

One potential solution to this is to approximate the history in the same way that DQN approximated the Q-values. Deep Recurrent Q Networks (DRQNs) are an extension to DQN that incorporate *recurrent* network layers. A recurrent network layer is any layer whose output is conditioned on previous inputs. What outputs are conditioned on is known as the *hidden state* of the layer and serves as an approximate memory.

Thus, by incorporating recurrent layers, a DRQN is able to infer the state from temporally extended inputs, and so its calculated $Q(o, a)$ can more accurately reflect sequences of events. The original DQN approximated this by providing input as a sequence of 4 frames*.

Related Work

The primary motivation behind this paper was Rainbow, a paper which combined several improvements to DQN*. Specifically, Rainbow applies:

- **Double Q-Learning**: Using a target network to reduce maximization bias*. Chooses the action with the policy network, but evaluates the chosen action with the target network. Loss:

$$(R_{t+1}+\gamma_{t+1}q_{\widehat{\theta}}(s_{t+1}, argmax_{a'}q_\theta(s_{t+1}, a')-q_\theta(s_t, a_t))^2 \quad (2)$$

- **Prioritized Replay**: Weighting experiences based on their difference from expected value*. Sample transitions with probability proportional to the last encountered *TD-error* at that transition.

$$p_t \propto |R_{t+1} + \gamma_{t+1} max_{a'} q_{\widehat{\theta}}(s_{t+1}, a') - q_\theta(s_t, a_t)|^\omega \quad (3)$$

- **Dueling Networks**: Dividing the output of a network into two streams corresponding to the value of a state and the relative advantage of taking an action*. The output is the value plus the best action minus the mean over all actions.
- **Multi-step Learning**: Calculating the return of multiple future steps rather than just the immediate reward and expected value of the next state*. The n-step return is:

$$R_t^{(n)} = \sum_{k=0}^{n-1} \gamma_t^{(k)} R_{t+k+1} \quad (4)$$

- **Distributional RL**: Maintaining a distribution of expected Q-values rather than a single one*. Goal is to minimize the KL divergence between the distribution d_t and the target distribution d_t'
- **Noisy Nets**: Replacing ϵ-greedy exploration (i.e., taking a random action with probability ϵ) with implicit exploration in the form of injected noise during training*. In place of a linear layer $y = b + Wx$:

$$y = (b + Wx) + (b_{noisy} \odot \epsilon^b + (W_{noisy} \odot \epsilon^w)x) \quad (5)$$

Rainbow and its ablations were tested extensively on the Atari environment in an attempt to determine which improvements made the biggest difference and explain the potential advantages of each. They concluded that prioritized replay and multi-step learning were the most impactful additions, but that the impacts of all the improvements were dependent on the environment chosen. Thus, it would be of interest to see how these improvements interact with more complex, partially-observable environments.

The other motivation was to leverage the VizDoom environment. VizDoom is a doom-based AI research platform specifically tailored to deep reinforcement learning*. As such, it has inspired both an annual AI deathmatch competition* and several papers. Among these is Chaplot and Lample's Arnold, an AI which combines DRQN navigation and DQN shooting and which won the 2017 competition. Arnold forms the primary basis for this paper's AI architecture.

Project Description

This project consisted of 3 main parts.

- Implement the algorithms from Rainbow for a DRQN-based agent
- Develop a scenario in VizDoom

- As in Rainbow, permute which algorithms are used for the agent and observe how the permutation affects training

For the first part, all networks are implemented in PyTorch. At first, I decided to re-use Arnold's architecture. As it had won a full deathmatch competition, it seemed that the architecture they used must have been a reasonable choice for the environment. I made one major change, though. Originally, Arnold took advantage of the game providing labels that flagged whether different entities were on the screen (e.g., flagging that an enemy is in view). They only used this during training, but it seemed an unfair advantage, so I removed it. The architecture is as follows:

- Each game frame is processed to 3x60x108 pixels. The authors noted that both color and a wide aspect ratio seemed to speed training.
- This input goes through two convolutional layers (with ReLU activation and batch normalization)
- Additionally, the game variables for health and ammo are processed through an embedding layer
- Both outputs are fed into a 1-layer LSTM with 512 hidden units
- The LSTM's output is mapped to actions through a linear layer
- For experiments using the dueling architecture, the final linear layer was replaced with two linear layers to serve as value and advantage streams. Their outputs were fused together as mentioned above.

Additionally, I adapted the code for prioritized replay and n-step learning from a different repository*. For this instantiation, I did not implement distributional RL due to the complexity.

After this modified architecture showed poor results, I wrote an entirely new version for a basic 1-frame DQN that heavily borrowed from SLM lab* and a series of pytorch tutorials*. In this case, I was able to adapt all 6 Rainbow algorithms.

When the new architecture also showed poor results, I extended it to a stacked-frame DQN similar to the original Atari paper.

For the second part, I had originally intended to define my own scenario. VizDoom provides a number of useful prebuilt scenarios, but I felt that most of them would be too simple. As it became clear that my agent was failing to learn, I turned back to the original scenarios. The scenario I tested my Arnold reimplementation on is called "Defend the Line", and the scenario I tested the others on is "Defend the Center". Both are described in the experiments section.

For the third part, I attempted to replicate the ablations of the rainbow paper. For each series of experiments, I did a run with the base architecture, with the fully augmented rainbow version, and then one run for each component to be removed in turn. I hypothesized that the VizDoom environment might have a different pattern of impact from these enhancements than the Atari games had.

Experiments

The scenarios are defined as follows:

- **Defend the Line**: The agent starts at the end of a hallway with limited ammunition. 3 melee and 3 ranged monsters spawn at the other end. These monsters can initially be killed by 1 shot, but as new ones spawn, they take progressively more damage.

  Reward: +1 for each kill

  Actions: Turn left, turn right, attack

- **Defend the Center**: The agent starts in the center of a room with limited ammunition. 3 melee monsters spawn randomly around the agent. All monsters can be killed with 1 shot.

  Reward: +1 for each kill

  Actions: Turn left, turn right, attack

In the LaTeX source for your paper, you **must** place the following lines as shown in the example in this subsection. This command set-up is for three authors. Add or subtract author and address lines as necessary, and uncomment the portions that apply to you. In most instances, this is all you need to do to format your paper in the Times font. The helvet package will cause Helvetica to be used for sans serif, and the courier package will cause Courier to be used for the typewriter font. These files are part of the PSNFSS2e package, which is freely available from many Internet sites (and is often part of a standard installation).

Leave the setcounter for section number depth commented out and set at 0 unless you want to add section numbers to your paper. If you do add section numbers, you must uncomment this line and change the number to 1 (for section numbers), or 2 (for section and subsection numbers). The style file will not work properly with numbering of subsubsections, so do not use a number higher than 2.

Conclusion

PDF files contain document summary information that enables us to create an Acrobat index (pdx) file, and also allows search engines to locate and present your paper more accurately. **Document Metadata for Author and Title are REQUIRED.**

If your paper includes illustrations that are not compatible with PDFTeX (such as .eps or .ps documents), you will need to convert them. The epstopdf package will usually work for eps files. You will need to convert your ps files to PDF however.

Important: Do not include *any* LaTeX code or nonascii characters (including accented characters) in the metadata. The data in the metadata must be completely plain ascii. It may not include slashes, accents, linebreaks, unicode, or any LaTeX commands. Type the title exactly as it appears on the paper (minus all formatting). Input the author names in the order in which they appear on the paper (minus all accents), separating each author by a comma. You may also include keywords in the Keywords field.

Preparing Your Paper

After the preamble above, you should prepare your paper as follows:

```
\begin{document}
\maketitle
...
\bibliography{Bibliography-File}
\bibliographystyle{aaai}
\end{document}
```

Incompatible Packages

The following packages are incompatible with aaai.sty and/or aaai.bst and must not be used (this list is not exhaustive — there are others as well):

- hyperref
- natbib
- geometry
- titlesec
- layout
- caption
- titlesec
- T1 fontenc package (install the CM super fonts package instead)

Illegal Commands

The following commands may not be used in your paper:

- \input
- \vspace (when used before or after a section or subsection)
- \addtolength
- \columnsep
- \top margin (or text height or addsidemargin or even side margin)

Paper Size, Margins, and Column Width

Papers must be formatted to print in two-column format on 8.5 x 11 inch US letter-sized paper. The margins must be exactly as follows:

- Top margin: .75 inches
- Left margin: .75 inches
- Right margin: .75 inches
- Bottom margin: 1.25 inches

The default paper size in most installations of LaTeX is A4. However, because we require that your electronic paper be formatted in US letter size, you will need to alter the default for this paper to US letter size. Assuming you are using the 2e version of LaTeX, you can do this by including the [letterpaper] option at the beginning of your file: \documentclass[letterpaper]article.

This command is usually sufficient to change the format. Sometimes, however, it may not work. Use PDFLaTeX and include \setlength{\pdfpagewidth}{8.5in} \setlength{\pdfpageheight}{11in} in your preamble.

Do not use the Geometry package to alter the page size. Use of this style file alters aaai.sty and will result in your paper being rejected.

Column Width and Margins. To ensure maximum readability, your paper must include two columns. Each column should be 3.3 inches wide (slightly more than 3.25 inches), with a .375 inch (.952 cm) gutter of white space between the two columns. The aaai.sty file will automatically create these columns for you.

Overlength Papers

If your paper is too long, turn on \frenchspacing, which will reduce the space after periods. Next, shrink the size of your graphics. Use \centering instead of \begin{center} in your figure environment. If these two methods don't work, you may minimally use the following. For floats (tables and figures), you may minimally reduce \floatsep, \textfloatsep, \abovecaptionskip, and \belowcaptionskip. For mathematical environments, you may minimally reduce \abovedisplayskip, \belowdisplayskip, and \arraycolsep. You may also alter the size of your bibliography by inserting \fontsize{9.5pt}{10.5pt} \selectfont right before the bibliography.

Commands that alter page layout are forbidden. These include \columnsep, \topmargin, \topskip, \textheight, \textwidth, \oddsidemargin, and \evensizemargin (this list is not exhaustive). If you alter page layout, you will be required to pay the page fee *plus* a reformatting fee. Other commands that are questionable and may cause your paper to be rejected include \parindent, and \parskip. Commands that alter the space between sections are also questionable. The title sec package is not allowed. Regardless of the above, if your paper is obviously "squeezed" it is not going to to be accepted. Before using every trick you know to make your paper a certain length, try reducing the size of your graphics or cutting text instead or (if allowed) paying the extra page charge. It will be cheaper in the long run.

Figures

Your paper must compile in PDFLaTeX. Consequently, all your figures must be .jpg, .png, or .pdf. You may not use the .gif (the resolution is too low), .ps, or .eps file format for your figures.

When you include your figures, you must crop them **outside** of LaTeX. The command \includegraphics*[clip=true, viewport 0 0 10 10]... might result in a PDF that looks great, but the image is **not really cropped.** The full image can reappear when page numbers are applied or color space is standardized.

Type Font and Size

Your paper must be formatted in Times Roman or Nimbus. We will not accept papers formatted using Computer Modern or Palatino or some other font as the text or heading typeface. Sans serif, when used, should be Courier. Use Symbol or Lucida or Computer Modern for *mathematics only.*

Do not use type 3 fonts for any portion of your paper, including graphics. Type 3 bitmapped fonts are designed for fixed resolution printers. Most print at 300 dpi even if the printer resolution is 1200 dpi or higher. They also often cause high resolution imagesetter devices and our PDF indexing software to crash. Consequently, AAAI will not accept electronic files containing obsolete type 3 fonts. Files containing those fonts (even in graphics) will be rejected.

Fortunately, there are effective workarounds that will prevent your file from embedding type 3 bitmapped fonts. The easiest workaround is to use the required times, helvet, and courier packages with LaTeX2e. (Note that papers formatted in this way will still use Computer Modern for the mathematics. To make the math look good, you'll either have to use Symbol or Lucida, or you will need to install type 1 Computer Modern fonts — for more on these fonts, see the section "Obtaining Type 1 Computer Modern.")

If you are unsure if your paper contains type 3 fonts, view the PDF in Acrobat Reader. The Properties/Fonts window will display the font name, font type, and encoding properties of all the fonts in the document. If you are unsure if your graphics contain type 3 fonts (and they are PostScript or encapsulated PostScript documents), create PDF versions of them, and consult the properties window in Acrobat Reader.

The default size for your type should be ten-point with twelve-point leading (line spacing). Start all pages (except the first) directly under the top margin. (See the next section for instructions on formatting the title page.) Indent ten points when beginning a new paragraph, unless the paragraph begins directly below a heading or subheading.

Obtaining Type 1 Computer Modern for LaTeX. If you use Computer Modern for the mathematics in your paper (you cannot use it for the text) you may need to download type 1 Computer fonts. They are available without charge from the American Mathematical Society: http://www.ams.org/tex/type1-fonts.html.

Title and Authors

Your title must appear in mixed case (nouns, pronouns, and verbs are capitalized) near the top of the first page, centered over both columns in sixteen-point bold type (twenty-four point leading). This style is called "mixed case." Author's names should appear below the title of the paper, centered in twelve-point type (with fifteen point leading), along with affiliation(s) and complete address(es) (including electronic mail address if available) in nine-point roman type (the twelve point leading). (If the title is long, or you have many authors, you may reduce the specified point sizes by up to two points.) You should begin the two-column format when you come to the abstract.

Formatting Author Information Author information can be set in a number of different styles, depending on the number of authors and the number of affiliations you need to display. For several authors from the same institution, use \and:

```
\author{Author 1 \and ... \and Author n\\
Address line \\ ... \\ Address line}
```

If the names do not fit well on one line use:

```
\author{Author 1}\\
{\bf Author 2}\\ ... \\ {\bf Author n}\\
Address line \\ ... \\ Address line}
```

For authors from different institutions, use \And:

```
\author{Author 1\\ Address line \\ ... \\ Address line
\And ... \And Author n\\
Address line\\ ... \\ Address line}
```

To start a separate "row" of authors, use \AND:

```
\author{Author 1\\ Address line \\ ... \\ Address line\\
\AND
Author 2 \\ Address line \\ ... \\ Address line\\
\And
Author 3 \\ Address line \\ ... \\ Address line\\
}
```

If the title and author information does not fit in the area allocated, place \setlength\titlebox{*height*} after the \documentclass line where {*height*} is something like 2.5in.

LaTeX Copyright Notice

The copyright notice automatically appears if you use aaai.sty. If you are creating a technical report, it is not necessary to include this notice. You may disable the copyright line using the \nocopyrightcommand. To change the entire text of the copyright slug, use: \copyrighttext {*text*}. Either of these must appear before \maketitle. Please be advised, however, that *if you disable or change the copyright line and transfer of copyright is required, your paper will not be published.*

Credits

Any credits to a sponsoring agency should appear in the acknowledgments section, unless the agency requires different placement. If it is necessary to include this information on the front page, use \thanks in either the \author or \title commands. For example:

```
\title{Very Important Results in AI\thanks{This work is supported by everybody.}}
```

Multiple \thanks commands can be given. Each will result in a separate footnote indication in the author or title with the corresponding text at the botton of the first column of the document. Note that the \thanks command is fragile. You will need to use \protect.

Please do not include \pubnote commands in your document.

Abstract

The abstract must be placed at the beginning of the first column, indented ten points from the left and right margins. The title Abstract should appear in ten-point bold type, centered above the body of the abstract. The abstract should be set in nine-point type with ten-point leading. This concise, one-paragraph summary should describe the general thesis and conclusion of your paper. A reader should be able to learn the purpose of the paper and the reason for its importance from the abstract. The abstract should be no more than two hundred words in length. (Authors who are submitting

short one- or two-page extended extracts should provide a short abstract of only a sentence or so.) **Do not include references in your abstract!**

Page Numbers

Do not **ever** print any page numbers on your paper.

Text

The main body of the paper must be formatted in ten-point with twelve-point leading (line spacing).

Citations

Citations within the text should include the author's last name and year, for example (Newell 1980). Append lowercase letters to the year in cases of ambiguity. Multiple authors should be treated as follows: (Feigenbaum and Engelmore 1988) or (Ford, Hayes, and Glymour 1992). In the case of four or more authors, list only the first author, followed by et al. (Ford et al. 1997).

Extracts

Long quotations and extracts should be indented ten points from the left and right margins.

> This is an example of an extract or quotation. Note the indent on both sides. Quotation marks are not necessary if you offset the text in a block like this, and properly identify and cite the quotation in the text.

Footnotes

Avoid footnotes as much as possible; they interrupt the reading of the text. When essential, they should be consecutively numbered throughout with superscript Arabic numbers. Footnotes should appear at the bottom of the page, separated from the text by a blank line space and a thin, half-point rule.

Headings and Sections

When necessary, headings should be used to separate major sections of your paper. Remember, you are writing a short paper, not a lengthy book! An overabundance of headings will tend to make your paper look more like an outline than a paper.

First-level heads should be twelve-point Times Roman bold type, mixed case (initial capitals followed by lower case on all words except articles, conjunctions, and prepositions, which should appear entirely in lower case), with fifteen-point leading, centered, with one blank line preceding them and three additional points of leading following them. Second-level headings should be eleven-point Times Roman bold type, mixed case, with thirteen-point leading, flush left, with one blank line preceding them and three additional points of leading following them. Do not skip a line between paragraphs. Third-level headings should be run in with the text, ten-point Times Roman bold type, mixed case, with twelve-point leading, flush left, with six points of additional space preceding them and no additional points of leading following them.

Section Numbers The use of section numbers in AAAI Press papers is optional. To use section numbers in LaTeX, uncomment the setcounter line in your document preamble and change the 0 to a 1 or 2. Section numbers should not be used in short poster papers.

Section Headings. Sections should be arranged and headed as follows:

Acknowledgments. The acknowledgments section, if included, appears after the main body of text and is headed "Acknowledgments." This section includes acknowledgments of help from associates and colleagues, credits to sponsoring agencies, financial support, and permission to publish. Please acknowledge other contributors, grant support, and so forth, in this section. Do not put acknowledgments in a footnote on the first page. If your grant agency requires acknowledgment of the grant on page 1, limit the footnote to the required statement, and put the remaining acknowledgments at the back. Please try to limit acknowledgments to no more than three sentences.

Appendices. Any appendices follow the acknowledgments, if included, or after the main body of text if no acknowledgments appear.

References The references section should be labeled "References" and should appear at the very end of the paper (don't end the paper with references, and then put a figure by itself on the last page). A sample list of references is given later on in these instructions. Please use a consistent format for references. Poorly prepared or sloppy references reflect badly on the quality of your paper and your research. Please prepare complete and accurate citations.

Illustrations and Figures

Figures, drawings, tables, and photographs should be placed throughout the paper near the place where they are first discussed. Do not group them together at the end of the paper. If placed at the top or bottom of the paper, illustrations may run across both columns. Figures must not invade the top, bottom, or side margin areas. Figures must be inserted using the \usepackage{graphicx}. Number figures sequentially, for example, figure 1, and so on.

The illustration number and caption should appear under the illustration. Labels, and other text in illustrations must be at least nine-point type.

Low-Resolution Bitmaps. You may not use low-resolution (such as 72 dpi) screen-dumps and GIF files—these files contain so few pixels that they are always blurry, and illegible when printed. If they are color, they will become an indecipherable mess when converted to black and white. This is always the case with gif files, which should never be used. The resolution of screen dumps can be increased by reducing the print size of the original file while retaining the same number of pixels. You can also enlarge files by manipulating them in software such as PhotoShop. Your figures should be a minimum of 266 dpi when incorporated into your document.

LaTeX Overflow. LaTeX users please beware: LaTeX will sometimes put portions of the figure or table or an equation in the margin. If this happens, you need to scale the figure or table down, or reformat the equation. Check your log file! You must fix any overflow into the margin (that means no overfull boxes in LaTeX). If you don't, the overflow text will simply be eliminated. **Nothing is permitted to intrude into the margins.**

Using Color. Your paper will be printed in black and white and grayscale. Consequently, because conversion to grayscale can cause undesirable effects (red changes to black, yellow can disappear, and so forth), we strongly suggest you avoid placing color figures in your document. Of course, any reference to color will be indecipherable to your reader.

Drawings. We suggest you use computer drawing software (such as Adobe Illustrator or, (if unavoidable), the drawing tools in Microsoft Word) to create your illustrations. Do not use Microsoft Publisher. These illustrations will look best if all line widths are uniform (half- to two-point in size), and you do not create labels over shaded areas. Shading should be 133 lines per inch if possible. Use Times Roman or Helvetica for all figure call-outs. **Do not use hairline width lines** — be sure that the stroke width of all lines is at least .5 pt. Zero point lines will print on a laser printer, but will completely disappear on the high-resolution devices used by our printers.

Photographs and Images. Photographs and other images should be in grayscale (color photographs will not reproduce well; for example, red tones will reproduce as black, yellow may turn to white, and so forth) and set to a minimum of 266 dpi. Do not prescreen images.

Resizing Graphics. Resize your graphics **before** you include them with LaTeX. You may **not** use trim or clip options as part of your \includgraphics command. Resize the media box of your PDF using a graphics program instead.

Fonts in Your Illustrations You must embed all fonts in your graphics before including them in your LaTeX document.

References

The aaai.sty file includes a set of definitions for use in formatting references with BibTeX. These definitions make the bibliography style fairly close to the one specified below. To use these definitions, you also need the BibTeX style file "aaai.bst," available in the author kit on the AAAI web site. Then, at the end of your paper but before \enddocument, you need to put the following lines:

```
\bibliographystyle{aaai} \bibliography{bibfile1,bibfile2,...}
```

The list of files in the \bibliography command should be the names of your BibTeX source files (that is, the .bib files referenced in your paper).

The following commands are available for your use in citing references:

\cite: Cites the given reference(s) with a full citation. This appears as "(Author Year)" for one reference, or "(Author Year; Author Year)" for multiple references.
\shortcite: Cites the given reference(s) with just the year. This appears as "(Year)" for one reference, or "(Year; Year)" for multiple references.
\citeauthor: Cites the given reference(s) with just the author name(s) and no parentheses.
\citeyear: Cites the given reference(s) with just the date(s) and no parentheses.

Warning: The aaai.sty file is incompatible with the hyperref and natbib packages. If you use either, your references will be garbled.

Formatted bibliographies should look like the following examples.

Book with Multiple Authors
Engelmore, R., and Morgan, A. eds. 1986. *Blackboard Systems.* Reading, Mass.: Addison-Wesley.

Journal Article
Robinson, A. L. 1980a. New Ways to Make Microcircuits Smaller. *Science* 208: 1019–1026.

Magazine Article
Hasling, D. W.; Clancey, W. J.; and Rennels, G. R. 1983. Strategic Explanations in Consultation. *The International Journal of Man-Machine Studies* 20(1): 3–19.

Proceedings Paper Published by a Society
Clancey, W. J. 1983b. Communication, Simulation, and Intelligent Agents: Implications of Personal Intelligent Machines for Medical Education. In Proceedings of the Eighth International Joint Conference on Artificial Intelligence, 556–560. Menlo Park, Calif.: International Joint Conferences on Artificial Intelligence, Inc.

Proceedings Paper Published by a Press or Publisher
Clancey, W. J. 1984. Classification Problem Solving. In *Proceedings of the Fourth National Conference on Artificial Intelligence,* 49–54. Menlo Park, Calif.: AAAI Press.

University Technical Report
Rice, J. 1986. Poligon: A System for Parallel Problem Solving, Technical Report, KSL-86-19, Dept. of Computer Science, Stanford Univ.

Dissertation or Thesis
Clancey, W. J. 1979b. Transfer of Rule-Based Expertise through a Tutorial Dialogue. Ph.D. diss., Dept. of Computer Science, Stanford Univ., Stanford, Calif.

Forthcoming Publication
Clancey, W. J. 1986a. The Engineering of Qualitative Models. Forthcoming.

Producing Reliable PDF Documents with LaTeX

Generally speaking, PDF files are platform independent and accessible to everyone. When creating a paper for a proceedings or publication in which many PDF documents must be merged and then printed on high-resolution PostScript RIPs, several requirements must be met that are not normally of concern. Thus to ensure that your paper will look like it does

when printed on your own machine, you must take several precautions:

- Use type 1 fonts (not type 3 fonts)
- Use only standard Times, Nimbus, and CMR font packages (not fonts like F3 or fonts with tildes in the names or fonts—other than Computer Modern—that are created for specific point sizes, like Times˜19) or fonts with strange combinations of numbers and letters
- Embed all fonts when producing the PDF
- Do not use the [T1]fontenc package (install the CM super fonts package instead)

Creating Output Using PDFLaTeX Is Required

By using the PDFTeX program instead of straight LaTeX or TeX, you will probably avoid the type 3 font problem altogether (unless you use a package that calls for metafont). PDFLaTeX enables you to create a PDF document directly from LaTeX source. The one requirement of this software is that all your graphics and images must be available in a format that PDFLaTeX understands (normally PDF).

PDFLaTeX's default is to create documents with type 1 fonts. If you find that it is not doing so in your case, it is likely that one or more fonts are missing from your system or are not in a path that is known to PDFLaTeX.

dvipdf Script Scripts such as dvipdf which ostensibly bypass the Postscript intermediary should not be used since they generally do not instruct dvips to use the config.pdf file.

dvipdfm Do not use this dvi-PDF conversion package if your document contains graphics (and we recommend you avoid it even if your document does not contain graphics).

Ghostscript

LaTeX users should not use GhostScript to create their PDFs.

Graphics

If you are still finding type 3 fonts in your PDF file, look at your graphics! LaTeX users should check all their imported graphics files as well for font problems.

Proofreading Your PDF

Please check all the pages of your PDF file. Is the page size A4? Are there any type 3, Identity-H, or CID fonts? Are all the fonts embedded? Are there any areas where equations or figures run into the margins? Did you include all your figures? Did you follow mixed case capitalization rules for your title? Did you include a copyright notice? Do any of the pages scroll slowly (because the graphics draw slowly on the page)? Are URLs underlined and in color? You will need to fix these common errors before submitting your file.

Improperly Formatted Files

In the past, AAAI has corrected improperly formatted files submitted by the authors. Unfortunately, this has become an increasingly burdensome expense that we can no longer absorb. Consequently, if your file is improperly formatted, it may not be possible to include your paper in the publication. If time allows, however, you will be notified via e-mail (with a copy to the program chair) of the problems with your file and given the option of correcting the file yourself (and paying a late fee) or asking that AAAI have the file corrected for you, for an additional fee. If you opt to correct the file yourself, please note that we cannot provide you with any additional advice beyond that given in your packet. Files that are not corrected after a second attempt will be withdrawn.

LaTeX 209 Warning

If you use LaTeX 209 we will not be able to publish your paper. Convert your paper to LaTeX2e.

Naming Your Electronic File

We request that you name your LaTeX source file with your last name (family name) so that it can easily be differentiated from other submissions. If you name your files with the name of the event or "aaai" or "paper" or "camera-ready" or some other generic or indecipherable name, you bear all risks of loss — it is extremely likely that your file may be overwritten.

Submitting Your Electronic Files to AAAI

Submitting your files to AAAI is a two-step process. It is explained fully in the author registration and submission instructions. Please consult this document for details on how to submit your paper.

Inquiries

If you have any questions about the preparation or submission of your paper as instructed in this document, please contact AAAI Press at the address given below. If you have technical questions about implementation of the aaai style file, please contact an expert at your site. We do not provide technical support for LaTeX or any other software package. To avoid problems, please keep your paper simple, and do not incorporate complicated macros and style files.

AAAI Press
2275 East Bayshore Road, Suite 160
Palo Alto, California 94303
Telephone: (650) 328-3123
E-mail: See the submission instructions for your particular conference or event.

Additional Resources

LaTeX is a difficult program to master. If you've used that software, and this document didn't help or some items were not explained clearly, we recommend you read Michael Shell's excellent document (testflow doc.txt V1.0a 2002/08/13) about obtaining correct PS/PDF output on LaTeX systems. (It was written for another purpose, but it has general application as well). It is available at www.ctan.org in the tex-archive.

Acknowledgments

AAAI is especially grateful to Peter Patel Schneider for his work in implementing the aaai.sty file, liberally using the ideas of other style hackers, including Barbara Beeton. We also acknowledge with thanks the work of George Ferguson for his guide to using the style and BibTeX files — which has been incorporated into this document — and Hans Guesgen, who provided several timely modifications, as well as the many others who have, from time to time, sent in suggestions on improvements to the AAAI style.

The preparation of the LaTeX and BibTeX files that implement these instructions was supported by Schlumberger Palo Alto Research, AT&T Bell Laboratories, Morgan Kaufmann Publishers, The Live Oak Press, LLC, and AAAI Press. Bibliography style changes were added by Sunil Issar. \pubnote was added by J. Scott Penberthy. George Ferguson added support for printing the AAAI copyright slug. Additional changes to aaai.sty and aaai.bst have been made by the AAAI staff.

Thank you for reading these instructions carefully. We look forward to receiving your electronic files!